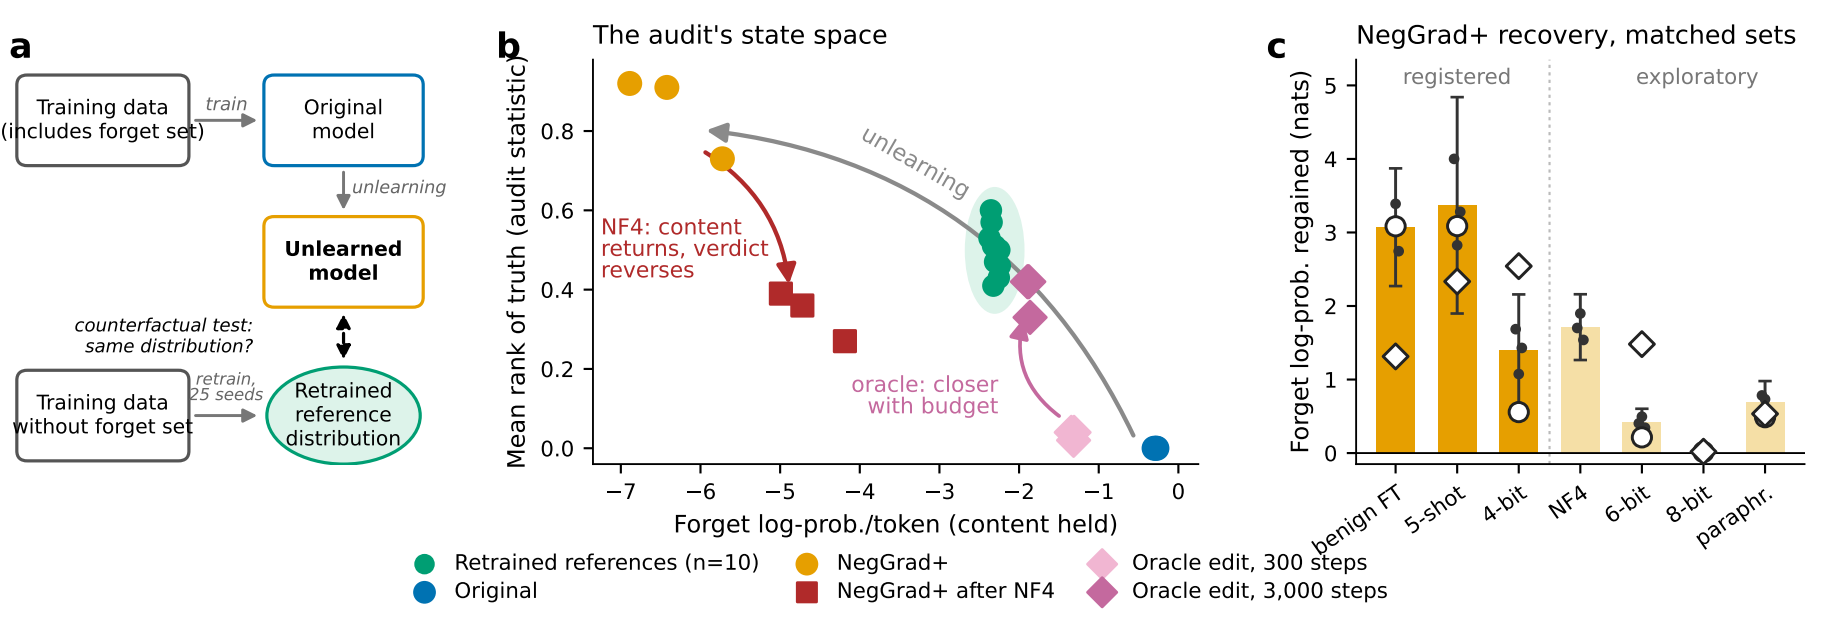

a
Training data (includes forget set)
train
Original model
unlearning
Unlearned model
counterfactual test: same distribution?
Training data without forget set
retrain, 25 seeds
Retrained reference distribution
b
The audit's state space
Mean rank of truth (audit statistic)
Forget log-prob./token (content held)
unlearning
NF4: content returns, verdict reverses
oracle: closer with budget
c
NegGrad+ recovery, matched sets
Forget log-prob. regained (nats)
registered
exploratory
benign FT
5-shot
4-bit
NF4
6-bit
8-bit
paraphr.
Retrained references (n=10)
Original
NegGrad+
NegGrad+ after NF4
Oracle edit, 300 steps
Oracle edit, 3,000 steps